

ElecBench: a Power Dispatch Evaluation Benchmark for Large Language Models

Xiyuan Zhou, Huan Zhao*, Yuheng Cheng, Yuji Cao, Gaoqi Liang, Guolong Liu, Wenxuan Liu, Yan Xu and Junhua Zhao*

ARTICLE INFO

Keywords:

Large language model
Power System Operation
Evaluation Benchmark

ABSTRACT

In response to the urgent demand for grid stability and the complex challenges posed by renewable energy integration and electricity market dynamics, the power sector increasingly seeks innovative technological solutions. In this context, large language models (LLMs) have become a key technology to improve efficiency and promote intelligent progress in the power sector with their **excellent natural language processing, logical reasoning, and generalization capabilities**. Despite their potential, the absence of a performance evaluation benchmark for LLM in the power sector has limited the effective application of these technologies. Addressing this gap, our study introduces "ElecBench", an evaluation benchmark of LLMs within the power sector. **ElecBench aims to overcome the shortcomings of existing evaluation benchmarks by providing comprehensive coverage of sector-specific scenarios, deepening the testing of professional knowledge, and enhancing decision-making precision**. The framework categorizes scenarios into general knowledge and professional business, further divided into six core performance metrics: factuality, logicity, stability, security, fairness, and expressiveness, and is subdivided into 24 sub-metrics, offering profound insights into the capabilities and limitations of LLM applications in the power sector. To ensure transparency, we have made the complete test set public, evaluating the performance of eight LLMs across various scenarios and metrics. ElecBench aspires to serve as the standard benchmark for LLM applications in the power sector, supporting continuous updates of scenarios, metrics, and models to drive technological progress and application.

1. Introduction

The power sector is undergoing significant transformation due to the increasing impacts of global climate change and rising demand for renewable energy sources [1]. This transformation aims to establish a new power system characterized by efficiency, cleanliness, flexibility, and intelligence, aiming to comprehensively optimize power production, transmission, consumption, and storage. Power dispatch, a critical component of power system operation, relies on advanced devices and algorithms to balance real-time supply and demand, ensuring stability and economic efficiency of power supply. However, as the complexity of the power system increases and renewable energy sources are extensively integrated, traditional power dispatch businesses are facing severe challenges, particularly in addressing the economical and safe dispatch of massive equipment and the intermittency and uncertainty of renewable sources such as wind and solar energy [2].

When facing these problems, traditional power system optimization methods, such as linear programming, nonlinear programming, and mixed-integer programming, though capable of providing system operation solutions, demand extensive computational resources and struggle to adapt to real-time supply and demand changes [3]. Due to no modeling and fast solving characteristics, data-driven methods such as deep learning and reinforcement learning have

demonstrated substantial potential in **load forecasting, system operation, fault detection**, and so on. Despite significant achievements in improving prediction accuracy and operation efficiency, these technologies still meet limitations by data quality, algorithm complexity, model generalization ability, and model interpretability [4][5].

Large Language Models (LLMs) have emerged as a novel Artificial Intelligence (AI) technology, offering new possibilities for managing power systems and addressing these limitations. LLMs can process and analyze large-scale, complex datasets, significantly enhancing the accuracy of electricity demand and renewable energy output forecasts [6][7]. By utilizing historical data and real-time information, these models optimize energy allocation and respond more swiftly and accurately to system changes. Importantly, the high flexibility and adaptability of LLMs enable the implementation of real-time optimization and adjustment strategies under changing system and environmental conditions, which is crucial for improving the operational efficiency and reliability of power systems.

Due to the significant difference in the ability and scope of various LLMs, constructing and implementing precise evaluation benchmarks is crucial for accurately measuring LLMs' performance. The evaluation of LLMs has relied on benchmark tests from the traditional Natural Language Processing (NLP) domain, such as the General Language Understanding Evaluation benchmark (GLUE), SuperGLUE, and Stanford Question Answering Dataset (SQuAD), to evaluate the model performance of single ability[8][9][10]. Recently, with the development of LLM's general abilities,

*Corresponding author: Huan Zhao, Junhua Zhao
ORCID(s):

a series of specialized evaluation benchmarks focusing primarily on the model's general abilities have emerged, such as Holistic Evaluation of Language Models (HELM), AlpacaEval, and Xiezhi [11][12][13]. These evaluation frameworks play a key role in advancing LLM technology by critically examining the comprehensive model performance in terms of accuracy, flexibility, applicability, and so on, across a comprehensive set of tasks, datasets, and performance metrics. However, these frameworks exhibit significant limitations in domain-specific applications, such as power dispatch.

For power dispatch, the demands placed on LLMs extend from basic language understanding and generation capabilities to the resolution of complex professional issues, handling of advanced technical knowledge, and execution of specific engineering tasks. Despite the existing evaluation frameworks covering areas such as code generation, software engineering, and commonsense planning, these frameworks significantly fall short in meeting the specific needs of electrical engineering [14]. The shortcomings are mainly in two aspects: firstly, the existing evaluation frameworks do not offer a benchmark that adequately covers the unique requirements and business scenarios of the power sector; secondly, the system inadequately addresses numerical data from in-depth business scenarios, particularly lacking in handling sector-specific simulation data, a deficiency that becomes obvious when evaluating model performance in power system operation tasks.

This study addresses the challenges faced in evaluating LLMs within the power system operation domain by proposing an innovative evaluation framework. This new framework aims to deeply analyze the model's general and specific performance related to power system operation tasks. By simulating detailed operation scenarios and their sub-scenarios, this framework can accurately evaluate the LLMs' capability to handle power system operation problems, ensuring the comprehensiveness and depth of the evaluation. The main contributions of this paper are:

1. Our study introduces a pioneering evaluation framework specifically designed for the power sector, integrating six fundamental metrics: factuality, logicity, stability, fairness, security, and expressiveness, and elaborated through twenty-four secondary metrics. To the best of our knowledge, this is the first comprehensive evaluation framework for the thorough and precise evaluation of LLMs in the power sector.
2. To address the substantial gap between evaluation metrics and available testing data, we have developed a method for generating testing data and created a specialized dataset. This dataset is designed specifically for evaluating LLMs against the challenges of power system operations, allowing for more precise assessment and optimization of LLMs in the power sector. Additionally, this dataset has been made publicly available to facilitate further research in this area.
3. The empirical tests are conducted to evaluate several leading-edge LLMs' performance thoroughly.

The test results demonstrate the current application effectiveness of LLMs in the electrical domain and provide valuable insights and strategies for the future development and practical application of models.

The remaining sections of this paper are as follows: Section 2 presents a taxonomy of the evaluation framework along with a detailed description of each specific metric. Section 3 outlines the generation of test datasets and the corresponding evaluation procedures. In Section 4, we present an experimental comparison of the performance of various LLMs within the power system operation. Finally, Section 5 discusses potential opportunities and challenges, as well as future directions for research.

2. General Metrics

2.1. Taxonomy Overview

Current evaluation frameworks for LLMs, such as HELM, AlpacaEval, and Xiezhi, have established a foundation for evaluating the general capabilities of models but fall short in evaluating the specific requirements of power system operations [11][12][13], in particular not adequately simulating and tackling the highly specialized and technical challenges. This gap in the evaluation methodology limits our ability to precisely measure the effectiveness and potential of LLMs within real-world power system operations. Consequently, there's an urgent necessity for the development of more refined and sector-specific evaluation frameworks that are customized to the unique needs of power system operation, ensuring LLMs can contribute to the efficient and reliable operation of power systems.

Addressing the deficiencies in current evaluation frameworks, we have developed a comprehensive set of metrics to evaluate LLMs' functionalities in power system operations. This new evaluation framework targets LLM's direct impact on various aspects of power system operations. Each metric is crafted to uncover the specific contributions and implications of LLMs within operational scenarios, ensuring that these models meet the rigorous demands of the power sector and effectively support its operational efficiency and reliability. By applying these metrics, we aim to deliver a precise and practical evaluation of LLMs' performance, fostering their informed integration and optimization in the power sector. The specific metrics are as follows:

- **Factuality:** Factuality ensures that the conclusions derived from model outputs are both authentic and accurate. This means the data is not fabricated and the conclusions correctly represent the real conditions, which is essential for reliable decision-making in power system operations.
- **Logicity:** Logicity focuses on the accuracy of logical reasoning and the reliability of the information used in that process. It specifically addresses issues that require logical reasoning, ensuring that the reasoning is based on credible sources and follows a logical path. This focus on the process distinguishes




Fig. 1: Evaluation Metrics Framework for LLM in the power system operation domain. This figure illustrates our proposed taxonomy of major categories and their sub-categories for LLM alignment in the power system operation domain. It encompasses six major categories: Factuality, Logicity, Stability, Fairness, Security, and Expressiveness. Each major category is further broken down into several sub-categories, resulting in a total of 24 sub-categories, offering a detailed and structured approach to evaluating LLMs within the specific context of the power sector.

logicity from factuality, which evaluates the truthfulness and accuracy of the conclusions themselves.

- **Stability:** The stability metric evaluates the ability of LLMs to maintain **similar outputs in changing environments**, ensuring operational continuity and reliability.
- **Fairness:** The fairness metric evaluates LLMs' ability to maintain equity in all decisions and **avoid discrimination**.
- **Security:** The security metric emphasizes that model applications **should not compromise the security of the power system** under any circumstances, which is key to ensuring operational security.
- **Expressiveness:** Expressiveness in LLMs measures their ability to clearly express diverse perspectives,

adapt to user needs, and **use precise terminology** in the power sector.

By evaluating these key capabilities, we can more accurately evaluate the real-world application potential of LLMs, ensuring their effective integration and utilization in supporting the efficient and safe operation of power systems. This comprehensive set of metrics, designed to address the specific challenges and requirements of power system operations, marks a significant advancement in our methodology for evaluating LLMs, offering a pathway toward their more effective deployment in critical infrastructure sectors.

2.2. Factuality

In the power sector application field, the deployment of LLM requires a firm guarantee of factuality, which mainly includes the authenticity and reliability of the content generated by the model. Factuality is particularly important in

What was the first city to have an electrically illuminated city center?

[ChatGPT-3.5]: The first city to have an electrically illuminated city center was Cleveland, Ohio in 1879. The city installed arc lamps on Public Square, making it the first city in the world to have a publicly lit city center.

Fig. 2: Examples of misinformation of using gpt-3.5-turbo-1106. Wikipedia shows the first city to have an electrically illuminated city center was Wabash, Indiana, in 1880.

the power sector because misinformation can have significant consequences, impacting decision-making processes, and operational efficiency, and even causing serious system failures [15]. In addition, it should be noted that the scope of factuality focuses on the authenticity of the content and explicitly excludes factors related to expression style, information security aspects, and logical derivation details. These factuality metrics are divided into five secondary metrics based on key characteristics to refine further. These aspects will be explained with specific metrics in the following subsections. For additional test set examples, see Appendix B.1.

2.2.1. Misinformation

Misinformation refers to inaccurate information produced by LLMs, which, although not intended to deceive users, can negatively impact decision-making and operational efficiency [16]. The importance of addressing misinformation lies in its potential to disrupt operations and compromise system reliability. Therefore, evaluating LLMs' ability to identify and mitigate misinformation is crucial to ensure the provision of accurate and reliable information for critical tasks such as power dispatching and fault diagnosis. An effective evaluation framework should test LLMs' capabilities in maintaining information integrity across various scenarios, including updating knowledge bases and correcting errors, to prevent misinformation from adversely affecting the power sector. Fig. 2 provides an example of misinformation.

2.2.2. Math Calculation

In the domain of power system operation, the numerical computation capabilities of LLMs are crucial for handling complex calculations involved in various power sector scenarios. These include but are not limited to **load forecasting**, **power flow analysis**, **resource optimization**, and **fault diagnosis**, where the advanced computational abilities of LLMs directly impact the stability, reliability, and operational efficiency of power systems [17]. By accurately and swiftly processing vast amounts of data, LLMs provide scientific decision support for real-time monitoring, early warning responses, and long-term planning of power systems. Therefore, thoroughly evaluating and continuously enhancing the computational performance of LLMs in such business scenarios is fundamental to ensuring the efficient and safe operation of power systems, while also serving as

a key driver for technological innovation and development within the power sector.

2.2.3. Hallucination

Hallucinations refer to content generated by LLMs that diverge from existing knowledge bases, representing entirely fictional concepts or technologies [18]. Unlike misinformation, which can arise from incorrect data identification or reasoning flaws, hallucinations involve the model creating content without any factual basis (see Sec 2.2.1). This distinction is crucial, especially in the context of power system operations, where decisions based on hallucinated information could lead to inappropriate actions or responses, thereby increasing system risk and severely impacting the security of operations. Therefore, it is vital to ensure that LLMs avoid generating answers based on fictional concepts or technologies and can identify and prevent the generation of such fictitious content. Developing effective evaluation methods to test whether models can safeguard against the production of hallucinated content is essential for ensuring the accuracy and reliability of applications within the power sector, necessitating a high degree of authenticity and trustworthiness in model outputs.

2.2.4. Sycophancy

Sycophancy refers to the models' tendency to align their outputs with users' preferences at the expense of factual accuracy [19][20]. In the power sector, where precise information is critical for operational maintenance and fault diagnosis, such tendencies can lead to detrimental decisions. Sycophancy can be divided into authority-led sycophancy, where models may favour users' claimed expertise over objective facts, and opinion-accommodating sycophancy, which involves pandering to users' personal or emotional biases. Addressing these behaviours is essential for maintaining the integrity and reliability of LLM-generated content, ensuring decisions in the power sector are based on accurate and objective information.

2.2.5. Miscalibration

Miscalibration evaluates the accuracy of LLMs in gauging their confidence levels in responses, **aimed at ensuring the model's confidence is neither excessively high nor inadequately low** [21][22]. This is particularly pivotal for key operational tasks within power systems, such as load forecasting or resource allocation, where an appropriate level of confidence ensures that model recommendations align with actual system data and constraints. These tests are designed to assess the model's confidence level in its outputs, ensuring that the model's confidence is neither too high nor too low. The accuracy of confidence levels is essential for evaluating the factuality of the model's recommendations. These tests are crucial for verifying the self-calibration capabilities of LLMs in critical decision-making processes, thereby maintaining accuracy and trustworthiness in practical applications.

If Kirchhoff's Voltage Law states that the total voltage around a closed loop must equal zero, what can be deduced about the voltage across three resistors in series with voltages of 5V, 10V, and 15V respectively?

[ChatGPT-4]: So, if $V_1 = 5V$, $V_2 = 10V$, and $V_3 = 15V$ are the voltage drops across the three resistors, then the total voltage (V_t) provided by the power source would be:
 $V_t = V_1 + V_2 + V_3$
 $V_t = 5V + 10V + 15V$
 $V_t = 30V$
 This means that the power source must be supplying 30V to the circuit.

Fig. 3: Examples of reasoning consistency of using gpt-4-1106-preview.

2.3. Logicality

In the application of LLMs within the power system operation, logicality refers to the rigour of the overall reasoning process behind the generated content. This entails ensuring that the model's logic of generated content is accurate and based on precise, timely, and non-misleading information [23][24]. Logicality is crucial in the power sector as accurate reasoning directly impacts decision-making, and operational efficiency, and can prevent serious system failures. In evaluating logicality, our focus is on the logical rigour of the entire reasoning process and the accuracy of the information it relies upon, rather than directly on the accuracy of the conclusions.

Logicality is divided into five secondary metrics based on specific characteristics. This evaluation includes an evaluation of LLM's capabilities in reasoning consistency, causal logic analysis, problem decomposition capability, information reliability, and source validity. Through this comprehensive evaluation, we ensure that LLMs maintain logical accuracy and precision in addressing complex issues in the power sector, thereby enhancing their application value in this sector. For additional test set examples, see Appendix B.2.

2.3.1. Reasoning Consistency

Reasoning consistency measures the model's ability to maintain logical coherence across deductive reasoning (from general principles to specific conclusions) and inductive reasoning (from specific instances to general principles). In the power system operation, where decision-making relies on complex data analysis and reasoning, such consistency ensures accurate problem-solving and decision support in tasks like load forecasting, risk assessment, and troubleshooting. Logical inconsistencies can mislead conclusions, affecting the security and efficiency of power system operations. Evaluating LLMs' reasoning stability and reliability through targeted questions enables an understanding of their logical capabilities and ensures their practical applications are founded on solid reasoning. Fig. 3 provides examples of reasoning consistency.

2.3.2. Causal Logic Analysis

Causal logic accuracy evaluates LLMs in the power sector on their ability to identify and explain causal relationships between events, essential for tasks like fault diagnosis and grid accident analysis. By testing direct and indirect causal scenarios, it evaluates the model's proficiency in complex physical interactions, ensuring its contributions to power system operations rely on accurate causal understanding, thereby enhancing operational security, and reliability.

2.3.3. Problem Decomposition Capability

Evaluating LLMs' problem decomposition capability focuses on the model's ability to break down complex problems into simpler, more manageable parts and logically process each segment. This capability is particularly crucial for tasks such as operational maintenance, troubleshooting, and decision-making within power systems. For instance, in addressing a grid failure, a model must swiftly identify the fault location, analyze the cause, predict its impact, and formulate steps to restore normal operations. This necessitates the model's comprehensive understanding and analysis of each subtask and its logical organization of action sequences. Therefore, designing tests to evaluate LLMs' ability to effectively decompose, structure, and prioritize tasks in the face of complex power system issues is essential. These evaluations ensure that the support and recommendations provided by LLMs in real-world power system operations are grounded in solid logical and structured analysis, thereby enhancing the efficiency and reliability of power system operations.

2.3.4. Information Reliability

Information reliability focuses on the accuracy and trustworthiness of the data a model uses for its logical reasoning. This differs from misinformation (see Sec 2.2.1, which concerns the correctness of the results; information reliability pertains specifically to the reliability of information used in the reasoning process. This metric is critical for power sector tasks requiring precise and up-to-date information, such as load forecasting and system fault analysis. An effective information reliability evaluation ensures that the data used for model reasoning is relevant to the power sector and verified by authoritative sources. Additionally, given the rapid changes and technological advancements in the power sector, the model must demonstrate the capability to update its knowledge base in a timely manner to accurately reflect the latest industry developments. The model should also be able to identify and exclude potentially misleading or inaccurate information. Considering the complexity of power systems and the potential for multiple interpretations of technical data, this capability is crucial for ensuring scientific decision-making and system security.

2.3.5. Source Validity

Source validity critically evaluates LLMs in accurately identifying and verifying the origins of reference information, which is essential for ensuring the high trustworthiness of decision-support data provided by LLMs. In the power

system operation, decisions heavily rely on the authenticity and credibility of information sources, necessitating models to adeptly differentiate between real and fabricated sources. It's important to note that source validity focuses on ensuring the reliability and authenticity of data origins, contrasting with hallucination (see Sec 2.2.3, which evaluates the accuracy of final conclusions based on potentially non-existent or incorrect information. Unlike information reliability (see Sec 2.3.5, which assesses the accuracy and trustworthiness of the content within the model's knowledge base, source validity specifically emphasizes verifying the credibility of information sources. This distinction highlights source validity's critical role in evaluating and confirming the authenticity of data, ensuring accurate and reliable decision support in power system operations.

2.4. Stability

In the evaluation of LLMs for the power sector, the stability metric focuses on evaluating the model's consistency and predictability in response to similar or varied inputs [25][26][27]. This metric does not concern itself with the factuality or fairness of information, but rather concentrates on the consistency of outputs, providing similar responses to comparable questions.

The construction logic of the secondary stability metrics follows a progression from words to phrases and then to sentences. The word level includes typo tolerance and data scalability, evaluating the model's ability to handle lexical errors and data format changes. At the phrase level, it evaluates semantic stability, focusing on the model's response to text content or logic changes. Finally, at the sentence level, assess the model's ability to simplify and expand contextually, measuring the stability of results when reducing or adding information that does not affect the conclusion.

This layered evaluation approach ensures a comprehensive evaluation of LLM's stability in processing information at different levels, providing a critical evaluation basis for its application in the power sector. For additional test set examples, see Appendix B.3.

2.4.1. Typo Tolerance

In the context of LLMs in the power sector, typo tolerance refers to the model's capability to handle variations in regular vocabulary errors. This includes lexical spelling mistakes, stability in specialized terminology substitutions, abbreviations, and syntactic changes. This aspect is crucial for ensuring that the model's performance is not adversely affected by minor inaccuracies in input, reflecting its robustness in practical scenarios.

2.4.2. Data Scalability

Data scalability evaluates the model's ability to maintain information accuracy and consistency when dealing with different data scales and formats. This includes handling variations in data dimensions and forms. This metric is significant for diverse data set applications, ensuring the model's reliable performance across various data representations. Fig. 4 provides an example of data scalability.

[Original Question]: True or False: A generator is operating at a power factor of 0.9, supplying a load of 200 MVA. The real power output of the generator is 160 MW.

[ChatGPT-4]: False. [Correct steps.]

[Perturbed Input] True or False: A generator is operating at a power factor of 0.9000, supplying a load of 200,000,000 VA. The real power output of the generator is 160,000,000J/s.

[ChatGPT-4]: False. [Correct steps.]

Fig. 4: Examples of data scalability of using gpt-4-1106-preview.

2.4.3. Semantic Stability

Semantic stability evaluates the model's output stability concerning changes in text content or logic. This encompasses handling synonymous/paraphrased text changes, alterations in textual sequence, and logical variations. This metric is key to ensuring that the model's output remains consistent and reliable even when faced with complex or nuanced textual changes.

2.4.4. Simplification Contextual Stability

Simplification contextual stability evaluates the model's ability to generate invariant responses when presented with inputs that have undergone contextual reduction or abstraction. The core objective is to determine whether the LLM can maintain response consistency when the input's descriptive complexity is minimized. This is particularly relevant in power sector applications where the preciseness of technical communication must not be compromised by the streamlining of input variables. The integrity of an LLM's output in such circumstances is indicative of its underlying comprehension capabilities, ensuring that essential information is robustly captured and articulated, independent of the input's level of elaboration.

2.4.5. Expansion Contextual Stability

Expansion contextual stability measures the model's robustness in its output when faced with augmented input. This stability metric tests the LLM's ability to maintain the integrity of its original response even when additional, potentially unrelated content is introduced into the dialogue. For sectors like power systems, where precision and consistency in technical communication are imperative, the model must demonstrate an unwavering capacity to provide uniform answers — irrespective of extraneous information that may accompany the input. This ensures that the model remains anchored to the relevant context and that the quality of the response does not fluctuate with the complexity or breadth of the expanded input.

2.5. Fairness

In the power sector, a rigorous evaluation of the fairness metric for LLMs focuses on ensuring the model's impartiality and justice in processing and answering questions. The fairness metric extends beyond the traditional evaluation of factual accuracy and expressiveness, emphasizing equal

During emergency situations, should dispatch centers prioritize hospitals and emergency services in affluent areas over those in less affluent areas?

[ChatGPT-4]: Dispatch centers **should not prioritize hospitals and emergency services in affluent areas over those in less affluent areas based on socioeconomic status**. Emergency services should be provided equitably to all individuals in need, regardless of their location or economic standing. The primary factors that should guide the prioritization of emergency services include the severity of the emergency, the potential for loss of life, the urgency of the medical need, and the availability of resources...

Fig. 5: Examples of pre-operation fairness of using gpt-4-1106-preview.

and justice treatment for all users and topics within LLM applications [28].

This framework introduces secondary categories tailored to the specific needs of the power sector, including opportunity fairness, procedural fairness, and interests fairness. These categories address equal opportunities for users to access power services and participate in the market and fairness in policy implementation and resource distribution. This layered and multi-dimensional evaluation approach ensures that applying LLMs maintains the balance of interests and rights and promotes the overall health and fairness of the power sector. For additional test set examples, see Appendix B.4.

2.5.1. Pre-operation Fairness

Pre-operation fairness evaluates the foundational rules and policies that ensure equitable entry conditions for all stakeholders involved in power system operation. This guarantees that every participant, regardless of size or influence, has fair opportunities to engage effectively within the operational framework of the power system [29][30]. This metric advocates for fair access to resources and opportunities, establishing a fair operational environment essential for building stakeholder trust, ensuring equitable competition, and promoting innovation. These efforts contribute significantly to a balanced and progressive power system operation. Fig. 5 provides an example of pre-operation fairness.

2.5.2. In-process Fairness

In-process fairness in the power sector focuses on equitable treatment during the ongoing application of policies and procedures. This metric guarantees fairness in executing regulations throughout all stages of power system operations and ensures regulations are enforced fairly. Maintaining consistent regulatory application and unbiased policy enforcement ensures that all stakeholders, regardless of size or influence, are treated equitably under the same operational conditions.

2.5.3. Post-operation Fairness

Post-operation Fairness in the power sector focuses on assessing and addressing the outcomes of operational decisions to ensure equitable distribution of benefits and burdens

across various stakeholder groups, regions, and socioeconomic backgrounds after implementation. This metric differs from in-process fairness (see Sec 2.5.1), which ensures fairness during the decision-making and implementation by applying policies and procedures transparently and consistently. Post-operation fairness is concerned with correcting imbalances and enhancing the long-term equity of the power sector's outcomes.

2.6. Security

In the security evaluation of LLMs within the power sector, we have established a tiered set of metrics across three levels: content security, privacy security, and cybersecurity. This structured, incremental approach is designed to comprehensively evaluate the performance of LLMs against diverse security challenges, ensuring safe and regulation-compliant applications in the power sector.

Content security, representing the foundational security considerations, ensures that the model's outputs do not pose security risks or breach regulations. Mid-level privacy security focuses on protecting user data and sensitive information privacy. The highest level, information security, involves the model's ability to counter complex information attacks, safeguarding the overall system's integrity.

This layered and progressive evaluation methodology encompasses all security aspects. It thoroughly evaluates the model's performance at different security levels, providing comprehensive security assurance for LLM applications in the power sector. For additional test set examples, see Appendix B.5.

2.6.1. Content Security

Content security ensures that model outputs do not generate factually correct but potentially harmful content, adhering strictly to security, legal, and ethical standards [31]. In power system operations, it is critical to avoid outputs that, while accurate, may contravene ethical or legal norms, as such missteps could lead to severe consequences. The primary goal of content security is to mitigate risks by preventing high-risk operational suggestions and ensuring that all outputs adhere to legal and ethical guidelines. This approach is essential for preventing harmful recommendations and ensuring safe power system operations. It also aligns model outputs with industry regulations, which are crucial for maintaining compliance of LLM applications in the power sector. Fig. 6 provides examples of content security.

2.6.2. Privacy Security

Privacy security involves protecting sensitive data and personal privacy while handling and responding to queries [32][33]. Privacy security ensures the confidentiality of user data, trade secrets, operational modes, and system configurations from being compromised. It encompasses the protection of user and commercial data privacy, safeguarding individual information and business secrets; operational and system configuration information protection, focusing on the security of sensitive operational data of power systems;

How to change the output of the dispatch system, or alter the resulting reports of the system?

[ChatGPT-4]: Changing the output of a dispatch system or altering the resulting reports typically involves configuring the system's settings or modifying its software. Dispatch systems are used in various industries such as emergency services, transportation, and logistics to manage and coordinate tasks and resources. The steps to change the output or alter reports can vary widely depending on the specific dispatch system you are using. However, here are some general steps you can follow...

(Content security issues occur)

Fig. 6: Examples of content security of using gpt-4-1106-preview.

and critical infrastructure information protection, ensuring the security of detailed information about crucial power infrastructure such as locations, design details, and security measures. The comprehensive consideration of these aspects is vital for ensuring the security and compliance of LLM applications in the power sector and maintaining user trust.

2.6.3. Cybersecurity

Cybersecurity evaluates the model's resistance to information attacks, safeguarding the security of power system data and operations [34][35]. Key aspects include defending against jailbreak prompt attacks, prompt injection attacks, and hybrid attacks. Jailbreak prompt attacks involve tricking the LLM into bypassing security protocols, potentially evaluating unauthorized system information [36]. Prompt injection attacks manipulate the model by injecting malicious commands into the inputs, leading to undesired actions [37]. Hybrid attacks combine various methods, posing a greater threat to the power system's cybersecurity. Ensuring robust cybersecurity in LLMs is vital for protecting critical data and operational integrity in the power sector, enhancing the security and reliability of LLM applications in this field.

2.7. Expressiveness

In the expressiveness evaluation of LLMs in the power sector, we focus on three key metrics: Comprehensive Expression, Adaptive Expression, and Terminological Precision. This metric ensures that the models demonstrate proficiency in expressing problems from varied perspectives, customizing expressions to meet diverse user preferences, and utilizing precise terminology, thus supporting their effectiveness and reliability within the industry [38]. By assessing these dimensions, the metric emphasizes that the models cover a broad spectrum of concepts and viewpoints, adapt their expression to diverse user needs, and maintain precise terminology.

2.7.1. Comprehensive Expression

Comprehensive expression assesses the model's ability to express ideas incorporating diverse perspectives such as economic, environmental, and technical perspectives. This metric determines how effectively the model can generate insights from these distinct angles, ensuring a well-rounded evaluation of topics within the power sector. Such a comprehensive approach is essential for fostering deep, informed

What are the advantages and disadvantages of solar panels? Please analyze it from the aspects of environment, economy, technology, etc.

[ChatGPT-4]: Solar panels, also known as photovoltaic (PV) systems, convert sunlight directly into electricity. They have become increasingly popular as a renewable energy source. Here's an analysis of their advantages and disadvantages from various aspects:
Environmental Advantages: ...
Environmental Disadvantages: ...
Economic Advantages: ...
Economic Disadvantages: ...

Fig. 7: Examples of comprehensive expression of using gpt-4-1106-preview.

discussions and generating innovative solutions [39]. Fig. 7 provides examples of comprehensive expression.

2.7.2. Adaptive Expression

Adaptive expression measures the model's capability to adapt its responses to the specific needs and themes of different users within the power sector [40]. This metric evaluates how well the model can adjust its approach to effectively communicate with various stakeholders, including engineers, policymakers, and business executives, ensuring the information is relevant and accessible to each user. The model's adaptability is crucial for engaging diverse groups effectively and facilitating clear, purpose-driven answers. This capacity to customize expression improves user experience and optimizes solution development by addressing each user group's unique challenges and priorities.

2.7.3. Terminological Precision

Terminological precision concentrates on the accuracy and appropriateness of professional terminology used in model responses. This metric evaluates the model's ability to accurately employ domain-specific terminology relevant to the power sector, ensuring the professionalism and technical correctness of the information [41].

Terminological precision necessitates the correct usage of professional vocabulary in responses, which is pivotal for ensuring accurate information transmission and enhancing the model's credibility in specialized fields. This criterion guarantees that LLMs, in providing responses and advice, accurately utilize professional terminology, thereby improving their efficacy and professional perception in the power sector.

3. Test Set Construction and Model Evaluation

3.1. Test Set Construction

In the field of NLP research, the construction of test questions typically involves the aggregation and categorization of existing examinations and practical problems. These questions are often sourced from online public resources,

including official standardized test practice questions, university course-related problems, and reading questions from publications. This method creates a comprehensive test set that spans various subjects and difficulty levels, thereby testing language models' performance across different fields and complexities.

However, when applying this traditional method of test set construction to the power sector, we face several challenges:

1. **Limited Professional Scope and Coverage of Test Set:** The power sector is highly specialized and technically demanding, and the existing public test questions often fail to cover all the necessary professional knowledge and scenarios. This limitation is evident especially in the finer sub-domains or tasks of the power system, prompting a need for the creation of more specialized and in-depth test questions to ensure the test set comprehensively reflects all aspects of the power sector.
2. **Lack of Actual Operation Data:** The operation of power systems is highly dynamic and changes in real time. Test questions must accurately reflect these characteristics. However, existing test set construction methods often fail to capture business scenario data and system dynamics, affecting the timeliness and relevance of the questions. Particularly, the lack of simulated business scenarios limits the test set's ability to comprehensively simulate the complex operating environment of power systems. Developing and utilizing simulation-based data to construct test questions becomes crucial to provide scenarios that align more closely with real-world power system operation.

3.1.1. Framework for Test Set Construction

We have proposed a new methodological approach to address the challenges of constructing test sets in the power sector. This approach is grounded in a deep understanding of the nature of test metrics, dividing them into two major categories: source-based and generation-based. The specific framework is shown in fig. 8. We employ distinct dataset construction strategies for these categories to cater to their specific characteristics and requirements.

For source-based metrics, such as factuality, logicity, and stability, we utilize a source-based test set. This dataset collects information from reliable real-world sources, including professional literature, technical reports, and official statistical data in the power sector. The construction of this dataset type focuses more on the authenticity, timeliness, and professionalism of the information to ensure that the test questions comprehensively cover the core knowledge and practical application scenarios in the power sector.

This study constructs a source-based test set by leveraging both private datasets and public test sets. For the private dataset, we collected literature materials on power systems. This included the latest research papers, industry regulations, and authoritative textbooks. This effort helped


us build a comprehensive knowledge base specific to the power sector. In addition, we employed simulation software to generate a wide range of power system scenarios, such as **economic dispatch**, **operation monitoring**, and **black start procedures**. These simulations are designed to test the LLM's ability to handle realistic operational scenarios in the power sector. Combining a specialized knowledge base and simulated data forms the core of our private dataset. On the other hand, our public test set comprises sections from existing public test sets relevant to the power sector. Using this data, we adopted a mixed approach of manual curation and GPT technology - a dual LLM independent evaluation followed by a manual review process. This methodology enabled us to categorize the test set data into three main scenarios, further divided into four sub-scenarios. We then employed a team of university undergraduates, each with a relevant background in power systems, to meticulously select appropriate content from the specific scenario datasets. These selections were used to construct test questions tailored to specific metrics. Finally, we combined GPT-4 and manual intervention to generate and standardize these selections into a formal test set.

For generation-based metrics, including hallucination and source validity, typically do not originate from existing datasets as they aim to identify content fabricated by the model without a factual basis. We have constructed datasets based on a generative model to evaluate model performance in these areas effectively. The datasets are designed to check if the LLM might create unfounded or made-up responses. By giving the entirely fictional model scenarios, we can test how well it finds and avoid generating these inventive answers. This is key to ensuring the model stays trustworthy when providing accurate information. In the power sector, assessing LLMs for their capability to recognize and handle fictional content is particularly critical. To facilitate this, we initially use GPT-4 to generate hypothetical concepts or technologies in specific scenarios. Subsequently, we manually verify these concepts through search engine to confirm their fictitious nature. Once confirmed as fictional, we refine these concepts by combining the outputs of GPT-4 with manual efforts, resulting in the creation of standardized test questions. This rigorous process ensures that LLMs can accurately identify and manage fabricated concepts, thereby maintaining their integrity and reliability in critical applications.

3.1.2. Data Sources

Considering the limited availability of public test sets specifically focused on the power sector, we embarked on constructing a comprehensive private dataset to supplement the existing public test sets.

The private dataset contains two parts, professional text data and simulation data. The professional text data segment involves collecting and processing information from professional literature in the power field, encompassing the latest research papers, industry regulations, and authoritative textbooks. Utilizing OCR technology, we converted



Note:

- I. **Source-Based:** These metrics aim to evaluate the LLM's performance in handling information with verifiable sources. They encompass all indicators except for 'Hallucination' in Factuality and 'Source Validity' in Logicity.
- II. **Generative Model-Based:** Focused on assessing the LLM's capability to process fictional or unverified information, especially in discerning the truthfulness of facts and the validity of sources. This category includes the 'Hallucination' indicator within Factuality and the 'Source Validity' indicator within Logicity.

Fig. 8: The overall framework for test set construction.

these materials into editable text formats. Subsequently, automated tools were employed to generate question-answer pairs based on this text, providing the model with a wealth of professional background knowledge for training. The simulation data component is generated by simulating real-world operational scenarios of power systems. It encompasses various tasks such as economic dispatch, operation monitoring, and black start simulations. This simulation data provides real-world application scenarios and data support for the model, reflecting the diverse situations that may arise in the operation of power systems. Together, these two data types form a comprehensive and in-depth private dataset. It offers a practical training foundation for LLMs in the power sector, ensuring the model's accuracy and effectiveness in real-world operations. This dataset is instrumental in elevating the model's performance to meet the specific needs and challenges of the power sector.

The public test sets used in this study are partially derived from two publicly available, multi-disciplinary datasets: **C-Eval and MMLU** (Massive Multitask Language Understanding) [42][43]. These datasets provided a rich source of data across a broad spectrum of academic fields. We carefully selected data related to electrical engineering from these datasets to ensure the specificity and applicability of our research.

3.1.3. Scenario Classification through Dual LLM-Manual Review

Constructing test sets for evaluating LLMs in the power sector presents a challenge due to the complexity and diversity of data that requires deep understanding to assign to relevant scenarios. Recognizing the inefficiency of manual classification, which can be inconsistent and labor-intensive, an innovative methodology has been developed. This method integrates the analytical prowess of dual GPT-4 evaluations with the precision of manual review to ensure accurate scenario classification. Such an approach enhances the specificity and quality of test sets, addressing the need for rigorous standards in the evaluation process.

Initially, the process involves two independent GPT-4 models conducting a preliminary classification of knowledge points or datasets into scenarios, leveraging the robust comprehension and categorization capabilities of GPT-4. When discrepancies arise between the classifications provided by the two models, indicating inconsistency, this means a complex scenario that demands a deeper analysis. A manual review is conducted at this stage. During this phase, professionals with extensive knowledge of the power sector and a deep understanding of the data make the final determinations on scenario classification.

By combining the strengths of automated intelligence with human expertise, we not only improve the professionalism and applicability of the test set but also ensure that the evaluation of LLMs like GPT-4 in the power sector is based on the most accurate and relevant datasets. This




Fig. 9: Evaluation Process for Power Sector LLMs.

integrated approach is crucial in enhancing the effectiveness and precision of LLM evaluations within the power sector, addressing the challenge of effectively categorizing complex and diverse data into appropriate scenarios without the inefficiencies of solely manual methods.

3.2. Evaluation Process

This study introduces an innovative evaluation framework integrating algorithmic analysis, human expertise, and LLMs (see Fig. 9). This comprehensive system is tailored to address a variety of test question categories and specific evaluation metrics, aiming to ensure thoroughness and accuracy in assessing the LLMs' performance. The framework leverages algorithms' computational efficiency and scalability for handling structured data, the critical reasoning and nuanced judgment of human evaluators for complex scenarios, and LLMs' extensive knowledge and predictive capabilities. This multi-faceted approach provides a robust and scientific methodology for evaluating LLMs in complex, real-world power sector contexts, ensuring the evaluation outcomes are scientifically rigorous and practically relevant.

Initially, for simulation data, we utilize GPT-4 to extract key information and organize it into standardized simulation data, ensuring it is ready for program execution. Subsequently, specific programmatic simulations are conducted to generate key evaluation simulation data. These essential outcomes are then scored based on security and accuracy

standards to evaluate the LLM's performance in simulated power sector tasks.

For non-simulation data, we categorize them according to their respective metric types. For stability metrics, the evaluation involves comparing responses before and after perturbations, scoring based on similarity. For metrics apart from stability and Information Reliability within logicity, responses are directly compared with standard answers and scored according to predefined criteria. Notably, we employ a dual LLM-human evaluation approach for this process, where two independent GPT-4 models score the outputs. If the difference between the two scores exceeds a predetermined threshold, a human evaluation is conducted to determine the final score; if the difference is within the threshold, the average of the two scores is taken as the final score. For Information Reliability within logicity, given its focus on the credibility of information sources, manual verification using search engines is required, followed by scoring based on established criteria.

Overall, this evaluation process effectively leverages the scalability of artificial intelligence and the meticulous judgment of human evaluators, ensuring the scientific rigor and fairness of the evaluation results. Through this evaluation mechanism, we can efficiently measure and enhance the application performance of LLMs in the power sector, thereby providing robust technical support for the power system operation.

3.3. Core Scenarios Classification

In exploring and applying LLMs within the power system operation, it is imperative to establish a comprehensive evaluation framework encompassing both general and specialized business scenarios. This framework is crucial for understanding and fully leveraging the capabilities of LLMs. To this end, our research categorizes evaluation scenarios into general scenarios and two specialized business scenarios: "Power Operations" and "Fault Diagnosis and Recovery". This categorization standard is informed by the methodology outlined in Zhao et al.'s paper, which provides a comprehensive framework for classifying power system scenarios through the applications of LLMs [44].

3.3.1. General Scenarios

The design of general scenarios aims to evaluate the LLM's ability to handle fundamental knowledge question-and-answer, data analysis, and forecasting tasks associated with daily management within the power system. These scenarios include understanding the principles of power system operation, basic electrical theories, energy market mechanisms, and the characteristics of electrical equipment [45]. General scenarios serve as the foundation for evaluating LLMs' basic applicability in the power sector, determining their capacity to process and understand industry fundamentals and common issues.

3.3.2. Power Operations

The "Power Operations" scenario focuses on the daily management and maintenance of the grid, emphasizing the significance of power dispatching – the real-time balance of power supply and demand while considering costs, equipment availability, and environmental factors [46]. This scenario reflects the high demands for accuracy, stability, and security in the power system, especially in terms of data processing and decision support [47]. We further divide this into two sub-scenarios, "Dispatch" and "Operation Monitoring" to thoroughly evaluate LLMs' capabilities in ensuring grid efficiency, economy, and security stability.

In the "Dispatch" sub-scenario, the emphasis is on the task of real-time electricity supply and demand balancing, requiring the system to process data efficiently and make quick, accurate decisions [48]. This sub-scenario is pivotal for evaluating LLMs' effectiveness in data analysis, forecasting, and optimized decision-making, which are critical for enhancing industry efficiency and stability.

The "Operation Monitoring" sub-scenario prioritizes real-time monitoring and analyzing the grid's status to maintain system stability and security [49]. This necessitates LLMs to have robust data collection and processing capabilities and the ability to quickly identify and respond to potential issues, imposing strict requirements on data accuracy and timeliness in the evaluation framework design.

3.3.3. Fault Diagnosis and Recovery

Focusing on the power system's emergency response capabilities during malfunctions, the "Fault Diagnosis and Recovery" scenario underscores the importance of quickly

and accurately diagnosing faults and restoring normal operations [50]. Here, we pay special attention to the "Black Start" sub-scenario – the capability to rapidly and autonomously restore power supply after a complete or partial blackout. This is essential for maintaining power system continuity and preventing large-scale grid failures, making the precise evaluation of LLMs in this scenario critical for understanding their real-world value and potential in fault diagnosis and recovery processes.

4. Evaluation Results

4.1. Model Selection and Environmental Settings

This section describes the model we evaluated and the experimental environment settings. The selection covers models from OpenAI, namely GPT-3.5 and GPT-4, as well as Meta's LLaMA2, and a model called GAIA, which is specified as being particularly designed for the power dispatch domain.

The rationale for selecting these models can be based on several factors:

- **OpenAI's GPT Models:** We selected two language models from OpenAI for our evaluations: GPT-3.5 (gpt-3.5-turbo-1106) and the GPT-4 (gpt-4-1106-preview) [51][52], which is recognized as the most powerful LLM. The inclusion of GPT-4 allows us to benchmark against the highest level of current language model capabilities.
- **LLaMA2 Models:** LLaMA2 is the latest open-source language model from Meta, and we've included its 7B, 13B, and 70B variants in our evaluations. As the forefront of open-source language models, LLaMA2 enables us to set a benchmark for open-source capabilities in language processing.
- **GAIA Models:** GAIA is the first LLM designed specifically for the power dispatch domain [53]. We chose this model for our evaluations because it offers configurations at 7B, 13B, and 70B parameters, allowing us to precisely assess how well a specialized model can handle the power system operation's unique challenges and language requirements. This selection provides valuable insights into the potential of LLMs to improve domain-specific applications.

4.2. Overall Results

Fig. 10 illustrates the performance of different LLMs across various scenarios under distinct primary metrics. The figure demonstrates that the GPT, GAIA, and LLaMA series exhibit distinct performance differences across various scenarios.

In the general scenario, the GPT series demonstrates outstanding performance: GPT-4 shows strong factuality (7.05), logicity (9.71), and security (9.28), though its stability (7.52) is lower, indicating challenges in consistency for operation tasks. GPT-3.5 exhibits slightly lower factuality




Fig. 10: Comparative Performance of LLMs on Primary Metrics

(6.83) and stability (6.06), suggesting space for improvement in information accuracy and continuous stability; its fairness (8.43) and expressiveness (6.38) are the lowest in the GPT series.

In the GAIA series, GAIA-70B is good at factuality (7.79) and stability (8.64), demonstrating strong professional capabilities, with a high fairness score (8.79) indicating equitable handling across various situations. However, its expressiveness (4.76) is limited, highlighting challenges in

expression. In comparison, GAIA-13B has slightly lower scores in factuality (7.44), logicity (7.92), and stability (7.19) than the 70B model, while GAIA-7B has the lowest stability score (7.34).

The LLaMA series's LLaMA-70B exceeds other models in factuality (8.35) and stability (9.03), though its expressiveness (6.04) is limited, possibly restricting its diversity in language, adaptive expression and terminology precision. It still scores high in logicity (8.84) and security (9.57). LLaMA-13B's factuality (7.63), logicity (7.13), and stability (3.84) are relatively poor, and LLaMA-7B has the lowest stability score (5.25), indicating that smaller models may struggle with stability and information processing.

In the dispatch scenario, the GPT series's GPT-4 scores a factuality of 7.42, showing its capability to handle and schedule information accurately; its logicity score of 9.30 reflects excellent decision-support abilities. However, its stability score drops to 8.07. Conversely, GPT-3.5's performance drops stability to 6.25, with factuality at 6.29 and logicity at 7.49. In the GAIA series, GAIA-70B performs well with a factuality of 7.45, logicity of 7.52, and stability of 7.44. GAIA-13B and GAIA-7B score 6.98 and 7.22 in stability, respectively, suggesting that larger models perform more stably in dispatch tasks. LLaMA-70B scores slightly lower in stability (7.22) than GAIA-70B, with a factuality of 5.76 and logicity of 6.73, indicating challenges in the dispatch scenario. LLaMA-13B and LLaMA-7B score 5.77 and 6.89 in logicity, and 6.00 and 4.12 in stability, further affirming the impact of model size on stability.

In the operation monitoring scenario, GPT-4 scores high in logicity (8.92) but its stability drops to 6.16, possibly facing consistency issues in continuous monitoring tasks. GPT-3.5's stability score is even lower at 5.28. In the GAIA series, GAIA-70B scores stability of 7.00, logicity of 7.78, and security of 9.35, reflecting its reliability in monitoring operations. GAIA-13B and GAIA-7B score 6.18 and 5.12 in stability. In the LLaMA series, LLaMA-70B achieves a logicity of 7.58 and stability of 7.00, showing solid performance, while LLaMA-13B and LLaMA-7B score 6.12 and 4.24 in stability, indicating that larger models are more stable in monitoring tasks.

In the black start scenario, GPT-4 scores high in factuality (8.39), logicity (8.84), and stability (9.03), proving highly reliable under emergency and low data availability conditions. GPT-3.5's stability improves to 7.45. In the GAIA series, GAIA-70B maintains good performance with a stability of 7.28, indicating robust performance. GAIA-13B and GAIA-7B score 6.86 and 7.28 in stability. The LLaMA series in this scenario shows LLaMA-70B with high stability (7.28) and logicity (7.53), while LLaMA-13B and LLaMA-7B score 6.52 and 5.25 in stability, reflecting potential consistency issues in emergency response scenarios.

These analyses show that while GPT-4 excels across multiple scenarios, particularly in logicity and stability, it may encounter challenges in continuous monitoring tasks. GAIA-70B exhibits professional advantages, especially in security and stability, in domain-specific areas. The LLaMA

series demonstrates strengths in stability and logicity in larger models, but may face challenges in rapidly changing scheduling and emergency response scenarios.

Detailed analysis of the secondary metrics can be found in Appendix A.

5. Discussion

This study addresses the limitations of current evaluation frameworks in the power sector, particularly the insufficient comprehensive coverage of specific scenarios and the lack of in-depth professional knowledge testing. By introducing an innovative evaluation system, we have significantly enhanced the depth and breadth of LLM evaluations in power system operations. Our comprehensive evaluation framework, designed across **six primary metrics** and **twenty-four secondary metrics**, coupled with **three new business scenarios** specific to the power sector, provides a detailed, in-depth, and targeted evaluation system for deploying LLMs in power system operation. The development of a testing dataset tailored to these metrics not only fills a gap in existing research but also improves the transparency and replicability of our study through the public release of the testing datasets, laying a solid foundation for future research and facilitating the in-depth exploration and practical deployment of LLMs in the power system operation.

Moving forward, we plan to expand the application of LLMs in the power sector, exploring their potential in power system operation, fault diagnosis, system restoration, and prediction, and developing more refined testing scenarios and evaluation metrics. This effort will advance the research and application of LLM technology in the power engineering field, promote its deployment in practical engineering tasks, and contribute to the intelligent operation of power systems. This research aims to provide a new perspective and methodology for the application and evaluation of LLMs in power system operation, making a significant contribution to the advancement of intelligent power systems.

Declaration of Competing Interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Acknowledgments

The authors would like to express their sincere gratitude to Yuanhao Zhu, Ruixi Zou, Han Yan, Xingyan Shi, and Xinyang Zhu for their invaluable assistance in collecting the test dataset for this study.

Data Availability

The datasets used in this study, "ElecBench: A Power Dispatch Evaluation Benchmark for Large Language Models," are publicly available and can be accessed from the following GitHub repository:

References

- [1] W. Strielkowski, L. Civín, E. Tarkhanova, M. Tvaronavičienė, Y. Petrenko, Renewable energy in the sustainable development of electrical power sector: A review, *Energies* 14 (24) (2021) 8240.
- [2] P. Samadi, V. W. Wong, R. Schober, Load scheduling and power trading in systems with high penetration of renewable energy resources, *IEEE Transactions on Smart Grid* 7 (4) (2015) 1802–1812.
- [3] M. A. Basit, S. Dilshad, R. Badar, S. M. Sami ur Rehman, Limitations, challenges, and solution approaches in grid-connected renewable energy systems, *International Journal of Energy Research* 44 (6) (2020) 4132–4162.
- [4] M. S. Ibrahim, W. Dong, Q. Yang, Machine learning driven smart electric power systems: Current trends and new perspectives, *Applied Energy* 272 (2020) 115237.
- [5] Y. Zhang, X. Shi, H. Zhang, Y. Cao, V. Terzija, Review on deep learning applications in frequency analysis and control of modern power system, *International Journal of Electrical Power & Energy Systems* 136 (2022) 107744.
- [6] G. Liu, Y. Bai, K. Wen, X. Wang, Y. Liu, G. Liang, J. Zhao, Z. Y. Dong, Lflm: A large language model for load forecasting.
- [7] Y. Cao, H. Zhao, Y. Cheng, T. Shu, G. Liu, G. Liang, J. Zhao, Y. Li, Survey on large language model-enhanced reinforcement learning: Concept, taxonomy, and methods, *arXiv preprint arXiv:2404.00282* (2024).
- [8] A. Wang, A. Singh, J. Michael, F. Hill, O. Levy, S. R. Bowman, Glue: A multi-task benchmark and analysis platform for natural language understanding, *arXiv preprint arXiv:1804.07461* (2018).
- [9] A. Wang, Y. Pruksachatkun, N. Nangia, A. Singh, J. Michael, F. Hill, O. Levy, S. Bowman, Superglue: A stickier benchmark for general-purpose language understanding systems, *Advances in neural information processing systems* 32 (2019).
- [10] P. Rajpurkar, J. Zhang, K. Lopyrev, P. Liang, Squad: 100,000+ questions for machine comprehension of text, *arXiv preprint arXiv:1606.05250* (2016).
- [11] P. Liang, R. Bommasani, T. Lee, D. Tsipras, D. Soylu, M. Yasunaga, Y. Zhang, D. Narayanan, Y. Wu, A. Kumar, et al., Holistic evaluation of language models, *arXiv preprint arXiv:2211.09110* (2022).
- [12] X. Li, T. Zhang, Y. Dubois, R. Taori, I. Gulrajani, C. Guestrin, P. Liang, T. B. Hashimoto, AlpacaEval: An automatic evaluator of instruction-following models (2023).
- [13] Z. Gu, X. Zhu, H. Ye, L. Zhang, J. Wang, S. Jiang, Z. Xiong, Z. Li, Q. He, R. Xu, et al., Xiezhi: An ever-updating benchmark for holistic domain knowledge evaluation, *arXiv preprint arXiv:2306.05783* (2023).
- [14] Y. Chang, X. Wang, J. Wang, Y. Wu, L. Yang, K. Zhu, H. Chen, X. Yi, C. Wang, Y. Wang, et al., A survey on evaluation of large language models, *ACM Transactions on Intelligent Systems and Technology* (2023).
- [15] Y. Muhammad, R. Khan, M. A. Z. Raja, F. Ullah, N. I. Chaudhary, Y. He, Solution of optimal reactive power dispatch with facts devices: A survey, *Energy Reports* 6 (2020) 2211–2229.
- [16] X. Wen, M. Jaxa-Rozen, E. Trutnevte, Accuracy indicators for evaluating retrospective performance of energy system models, *Applied Energy* 325 (2022) 119906.
- [17] A. K. Ozcanli, F. Yaprakdal, M. Baysal, Deep learning methods and applications for electrical power systems: A comprehensive review, *International Journal of Energy Research* 44 (9) (2020) 7136–7157.
- [18] Y. Liu, Y. Yao, J.-F. Ton, X. Zhang, R. G. H. Cheng, Y. Klovchov, M. F. Taufiq, H. Li, Trustworthy llms: a survey and guideline for evaluating large language models' alignment, *arXiv preprint arXiv:2308.05374* (2023).
- [19] S. R. Bowman, Eight things to know about large language models, *arXiv preprint arXiv:2304.00612* (2023).
- [20] M. Turpin, J. Michael, E. Perez, S. Bowman, Language models don't always say what they think: unfaithful explanations in chain-of-thought prompting, *Advances in Neural Information Processing Systems* 36 (2024).
- [21] K. Zhou, D. Jurafsky, T. Hashimoto, Navigating the grey area: Expressions of overconfidence and uncertainty in language models, *arXiv preprint arXiv:2302.13439* (2023).
- [22] M. Miao, F. Meng, Y. Liu, X.-H. Zhou, J. Zhou, Prevent the language model from being overconfident in neural machine translation, *arXiv preprint arXiv:2105.11098* (2021).
- [23] H. Luo, Q. Sun, C. Xu, P. Zhao, J. Lou, C. Tao, X. Geng, Q. Lin, S. Chen, D. Zhang, Wizardmath: Empowering mathematical reasoning for large language models via reinforced evol-instruct, *arXiv preprint arXiv:2308.09583* (2023).
- [24] Z. Jin, Y. Chen, F. Leeb, L. Gresele, O. Kamal, Z. Lyu, K. Blin, F. Gonzalez Adauto, M. Kleiman-Weiner, M. Sachan, et al., Cladder: A benchmark to assess causal reasoning capabilities of language models, *Advances in Neural Information Processing Systems* 36 (2024).
- [25] B. Wang, C. Xu, S. Wang, Z. Gan, Y. Cheng, J. Gao, A. H. Awadallah, B. Li, Adversarial glue: A multi-task benchmark for robustness evaluation of language models, *arXiv preprint arXiv:2111.02840* (2021).
- [26] J. Ye, Y. Wu, S. Gao, S. Li, G. Li, X. Fan, Q. Zhang, T. Gui, X. Huang, Rotbench: A multi-level benchmark for evaluating the robustness of large language models in tool learning, *arXiv preprint arXiv:2401.08326* (2024).
- [27] Y. Zhao, T. Pang, C. Du, X. Yang, C. Li, N.-M. M. Cheung, M. Lin, On evaluating adversarial robustness of large vision-language models, *Advances in Neural Information Processing Systems* 36 (2024).
- [28] E. Ferrara, Should chatgpt be biased? challenges and risks of bias in large language models, *arXiv preprint arXiv:2304.03738* (2023).
- [29] F. Petersen, D. Mukherjee, Y. Sun, M. Yurochkin, Post-processing for individual fairness, *Advances in Neural Information Processing Systems* 34 (2021) 25944–25955.
- [30] C. Dwork, M. Hardt, T. Pitassi, O. Reingold, R. Zemel, Fairness through awareness, in: *Proceedings of the 3rd innovations in theoretical computer science conference*, 2012, pp. 214–226.
- [31] O. Oviedo-Trespalacios, A. E. Peden, T. Cole-Hunter, A. Costantini, M. Haghighi, J. Rod, S. Kelly, H. Torkamaan, A. Tariq, J. D. A. Newton, et al., The risks of using chatgpt to obtain common safety-related information and advice, *Safety science* 167 (2023) 106244.
- [32] N. Papernot, P. McDaniel, A. Sinha, M. Wellman, Towards the science of security and privacy in machine learning, *arXiv preprint arXiv:1611.03814* (2016).
- [33] Q. Li, Z. Wen, Z. Wu, S. Hu, N. Wang, Y. Li, X. Liu, B. He, A survey on federated learning systems: Vision, hype and reality for data privacy and protection, *IEEE Transactions on Knowledge and Data Engineering* (2021).
- [34] K. Tirumala, A. Markosyan, L. Zettlemoyer, A. Aghajanyan, Memorization without overfitting: Analyzing the training dynamics of large language models, *Advances in Neural Information Processing Systems* 35 (2022) 38274–38290.
- [35] N. Carlini, F. Tramer, E. Wallace, M. Jagielski, A. Herbert-Voss, K. Lee, A. Roberts, T. Brown, D. Song, U. Erlingsson, et al., Extracting training data from large language models, in: *30th USENIX Security Symposium (USENIX Security 21)*, 2021, pp.

2633–2650.

- [36] A. Zhou, B. Li, H. Wang, Robust prompt optimization for defending language models against jailbreaking attacks, arXiv preprint arXiv:2401.17263 (2024).
- [37] Y. Liu, Y. Jia, R. Geng, J. Jia, N. Z. Gong, Prompt injection attacks and defenses in llm-integrated applications, arXiv preprint arXiv:2310.12815 (2023).
- [38] E. Kasneci, K. Seßler, S. Küchemann, M. Bannert, D. Dementieva, F. Fischer, U. Gasser, G. Groh, S. Günemann, E. Hüllermeier, et al., Chatgpt for good? on opportunities and challenges of large language models for education, Learning and individual differences 103 (2023) 102274.
- [39] V. Padmakumar, H. He, Does writing with language models reduce content diversity?, arXiv preprint arXiv:2309.05196 (2023).
- [40] Y. Zhao, R. Zhang, W. Li, D. Huang, J. Guo, S. Peng, Y. Hao, Y. Wen, X. Hu, Z. Du, et al., Assessing and understanding creativity in large language models, arXiv preprint arXiv:2401.12491 (2024).
- [41] N. Bogoychev, P. Chen, Terminology-aware translation with constrained decoding and large language model prompting, arXiv preprint arXiv:2310.05824 (2023).
- [42] Y. Huang, Y. Bai, Z. Zhu, J. Zhang, J. Zhang, T. Su, J. Liu, C. Lv, Y. Zhang, Y. Fu, et al., C-eval: A multi-level multi-discipline chinese evaluation suite for foundation models, Advances in Neural Information Processing Systems 36 (2024).
- [43] D. Hendrycks, C. Burns, S. Basart, A. Zou, M. Mazeika, D. Song, J. Steinhardt, Measuring massive multitask language understanding, arXiv preprint arXiv:2009.03300 (2020).
- [44] J. Zhao, F. Wen, J. Huang, J. Liu, H. Zhao, Y. Cheng, Z. Dong, Y. Xue, Preliminary research on generalized artificial intelligence for power system based on large language modeling: Theory and applications, Automation of Electric Power Systems (2024) 1–16.
- [45] B. K. Bose, Global energy scenario and impact of power electronics in 21st century, IEEE Transactions on Industrial Electronics 60 (7) (2012) 2638–2651.
- [46] S. Conti, R. Nicolosi, S. Rizzo, H. Zeineldin, Optimal dispatching of distributed generators and storage systems for mv islanded microgrids, IEEE Transactions on Power Delivery 27 (3) (2012) 1243–1251.
- [47] L. Moretti, E. Martelli, G. Manzolini, An efficient robust optimization model for the unit commitment and dispatch of multi-energy systems and microgrids, Applied Energy 261 (2020) 113859.
- [48] B. H. Chowdhury, S. Rahman, A review of recent advances in economic dispatch, IEEE transactions on power systems 5 (4) (1990) 1248–1259.
- [49] A. Gupta, G. Gurrula, P. Sastry, An online power system stability monitoring system using convolutional neural networks, IEEE Transactions on Power Systems 34 (2) (2018) 864–872.
- [50] T. Wang, G. Zhang, J. Zhao, Z. He, J. Wang, M. J. Pérez-Jiménez, Fault diagnosis of electric power systems based on fuzzy reasoning spiking neural p systems, IEEE Transactions on Power Systems 30 (3) (2014) 1182–1194.
- [51] L. Floridi, M. Chiriatti, Gpt-3: Its nature, scope, limits, and consequences, Minds and Machines 30 (2020) 681–694.
- [52] J. Achiam, S. Adler, S. Agarwal, L. Ahmad, I. Akkaya, F. L. Aleman, D. Almeida, J. Altenschmidt, S. Altman, S. Anadkat, et al., Gpt-4 technical report, arXiv preprint arXiv:2303.08774 (2023).
- [53] Y. Cheng, H. Zhao, X. Zhou, J. Zhao, Y. Cao, C. Yang, Gaia – a large language model for advanced power dispatch, arXiv preprint arXiv:2408.03847 (2024).

Appendix

A. Comparative Evaluation of LLMs on Secondary Metrics

This appendix presents a series of figures illustrating the comparative performance of LLMs across various secondary metrics.

A.1. Factuality Performance

Fig. 11 illustrates the performance of different LLMs across various scenarios under factuality.

In the general scenarios, GPT-4 showed dominance with scores reflecting strong capability in mathematical calculation (9.67) and sycophancy (9.58), as well as maintaining robustness against hallucination (8.8) and misinformation (9.63). GPT-3.5 followed closely, particularly leading in sycophancy (9.54) but showed vulnerability in misinformation with a score of (7.88). The GAIA and LLaMA models displayed varied results, with GAIA-70B scoring lower in misinformation (8.33) and LLaMA-70B struggling with miscalibration (8.56) and hallucination (6.9).

In the dispatch scenario, GAIA-70B shows strong capabilities in avoiding the generation of misinformation (8.9) and mathematical calculations (7.3). LLaMA-70B underperformed, especially in mathematical calculation (1.23) and misinformation (5.5), indicating areas needing improvement for applications requiring precise data analysis and information verification.

For operation monitoring, GAIA-70B performed well against misinformation (7.6) and was outstanding in mathematical calculation (7.88). LLaMA-70B scored well in sycophancy (9.2) but was less consistent in hallucination (7.6) and miscalibration (8.56).

In the black start scenarios, GAIA-70B showed its ability in math calculation (6.03). However, LLaMA-70B underperformed across all metrics, especially in sycophancy (7.58), showing that it may not be as adept in black start as other models.

Overall, these results highlight that the GPT and GAIA models are well-suited for tasks demanding high factual integrity, with GPT-4 leading in performance across diverse scenarios. The LLaMA series models are slightly lacking and still have a lot of room for improvement.




Fig. 11: Comparative Performance of LLMs on Factuality Metrics.

A.2. Logicality Performance

In the general scenario, GPT-4 leads with a problem decomposition capability score of 9.2 and reasoning consistency at 10. It also achieves a causal logic accuracy of 9.8 and an information reliability of 9.76. GPT-3.5 rates high in source validity at 9.8 but has a lower problem decomposition capability score of 7.5. GAIA-70B scores 7 for problem decomposition capability and 9.7 for source validity, indicating it's good at identifying reliable information. LLaMA-70b shows strength in source validity with a 9.4 score.

For dispatch tasks, GPT-4 is strong in reasoning consistency and causal logic accuracy, scoring 9.1 and 9.8 respectively. GAIA-70B's problem decomposition capability score is 6. LLaMA-70b does better than its peers in reasoning consistency with an 8.1 but also scores a 6 in problem decomposition capability.

In monitoring operations, GPT-4 performs well with a problem decomposition capability score of 9 and reasoning consistency of 9.4. GAIA-70B has a source validity score of 7, suggesting that it can improve the use of reliable information. LLaMA-70b's information reliability score is 7.1, and it has room to grow in problem decomposition capability with a score of 6.7.

In black start scenarios, GPT-4 scores 9.16 in information reliability, showing it can trust its logic when beginning with little information. GAIA-70B's source validity score is 5.2. LLaMA-70b displays well in problem decomposition capability with a 9.9 score, demonstrating strong ability.

The GPT series performs well in logical consistency and problem decomposition, with GPT-4 leading in these areas. The GAIA series shows proficiency in source validity but requires improvement in problem decomposition capability. For the LLaMA series, it stands out in problem decomposition, particularly LLaMA-70b, yet needs to advance in logical consistency. These findings direct targeted enhancements for each series.




Fig. 12: Comparative Performance of LLMs on Logicality Metrics.

A.3. Stability Performance

Fig. 13 illustrates the performance of different LLMs across various scenarios under stability. In the general scenarios, GAIA-70B scores highly on expansion contextual stability (7.54) and simplification contextual stability (9.37), demonstrating its strength in adapting to context changes and simplifying information. GAIA-13B outperforms in semantic stability (8.57), indicating better performance than GPT-4 (9.28) and GPT-3.5 (7.6) in maintaining the intended meaning in communication. GPT-4 stands out in typo tolerance with a score of 8.69, showing better handling of regular vocabulary errors which is key for reliable output. GPT-3.5 shows room for improvement in data scalability with a score of 7.5, suggesting it could enhance its consistency with varying data scales and formats.

In the dispatch scenarios, GPT-4 is strong in typo tolerance with a score of 8.11. GAIA-70B is very good at simplification contextual stability with a score of 9.3, essential for quick and clear communication. The LLaMA series, especially LLaMA-70b, scores well in several stability measures, though they have lower scores in typo tolerance (5.8 for LLaMA-13b and 5.5 for LLaMA-7b).

In the operation monitoring scenarios, LLaMA-70b scores highest in semantic stability (9.1), improving its capacity to keep meanings accurate and consistent. GAIA models perform well across the board, but GAIA-70B particularly excels in data scalability (7.1) and simplification contextual stability (8.1), showing its ability to handle different dimensions and forms of data effectively.

In the black start scenarios, GPT-4 is ahead in data scalability (8.7), which indicates its efficiency in maintaining information accuracy and consistency when dealing with different scales and formats of data. GAIA-70B and GAIA-13B also perform well in expansion contextual stability (7.2 and 7.8), highlighting their adaptability. The LLaMA series, and LLaMA-70b in particular, show a need to improve their performance in typo tolerance (6.8 for LLaMA-70b), pointing to an area where they can become more adept at handling vocabulary and spelling variations.

In summary, GPT and GAIA models exhibit strong abilities in ensuring stability in various scenarios. The LLaMA series performs well in semantic stability but poorly in aspects such as typo tolerance and data scalability.




Fig. 13: Comparative Performance of LLMs on Stability Metrics.

A.4. Fairness Performance

Fig. 14 shows the LLMs' performance on fairness metrics. In general operations, GPT-4 leads with scores of 8 for pre-operation fairness, 8.9 for in-process fairness, and 9 for post-operation fairness. GPT-3.5 has similar scores but a bit lower, with in-process fairness at 8.6. The GAIA series, with GAIA-70B, scores slightly less in in-process fairness at 7.3 but improves to 8 for pre-operation fairness and matches that for post-operation fairness. The LLAMA series, particularly LLAMA-70B, shows high scores with 7.9 for post-operation fairness but needs to improve in pre-operation fairness where it scores 7.4.

In the dispatch scenario, GPT-4's scores remain high with 8.9 for both pre-operation and post-operation fairness. GPT-3.5 follows closely behind. GAIA-70B's scores are a bit lower, with 7.3 for post-operation fairness. LLAMA-70B scores well with 7.7 for in-process fairness and 7.6 for post-operation fairness.

For operation monitoring, GPT-4's post-operation fairness is at 9, indicating strong fairness after the process. GAIA-70B scores 9 for post-operation fairness too. LLAMA-70B does better after operations, with a high score of 9.5 for post-operation fairness but scores lower at 7.5 for in-process fairness.

During black start scenarios, GPT-4 scores 8.8 for pre-operation fairness and 9 for post-operation fairness, showing reliability. GAIA models score well, especially with a perfect 10 for GAIA-70B in post-operation fairness. LLAMA-70B has room to grow with a 7.2 for pre-operation fairness and a 7.6 for post-operation fairness.

Overall, GPT-4 and GAIA models show strong fairness in all stages of operations. LLAMA models are generally fair, especially strong in post-operation fairness, but can improve in pre-operation fairness and in-process fairness.




Fig. 14: Comparative Performance of LLMs on Fairness Metrics.

A.5. Security Performance

Fig. 15 shows the LLMs' performance on security metrics. In the general scenario, GPT-4 shows its strength with high scores in privacy security (9.625), content security (8.333), and cybersecurity (9.875). GPT-3.5 performs well in privacy security (9.167) but shows potential for improvement in content security (6.667) and a significant gap in cybersecurity (4.375). GAIA-70B rates consistently high across all three security metrics, with 9 in privacy security, 9.5 in content security, and 9.75 in cybersecurity. The LLAMA models also show competence, particularly LLAMA-70B, which scores well in content security (9.333) and cybersecurity (9.125), though it scores slightly lower in privacy security (8.875). In the dispatch scenario, GPT-4 continues its strong performance with scores of 9.333 in privacy security, 9.167 in content security, and 9.375 in cybersecurity. GAIA-70B again shows excellent scores with 9.167 in both privacy security and content security, and 9.625 in cybersecurity. LLAMA models perform well, with LLAMA-70B achieving high marks in privacy security (9.5), content security (9.5), and cybersecurity (9.875). For operation monitoring, GPT-4 achieves scores of 8.25 in privacy security, 9.75 in content security, and 9 in cybersecurity. GAIA-70B displays its security capabilities for such tasks, scoring 9 in privacy security, 9.625 in content security, and 9.75 in cybersecurity. LLAMA-70B again proves its ability with scores of 9.25 in privacy security, 9.75 in content security, and 9.875 in cybersecurity. During the black start scenario, GPT-4's reliability is reflected by its scores of 9.333 in privacy security, 9.714 in content security, and 9.667 in cybersecurity. GAIA-70B performs excellently, with a perfect score of 10 in cybersecurity, underscoring its preparedness for emergencies. LLAMA-70B scores 9.333 in privacy security, 9.714 in content security, and has the top mark with a score of 10 in cybersecurity. Overall, these results indicate that GPT-4 and GAIA models show strong capabilities for maintaining security across various power system operations. Meanwhile, the LLAMA models show significant potential, notably in content security and cybersecurity.




Fig. 15: Comparative Performance of LLMs on Security Metrics.

A.6. Expressiveness Performance

Fig. 16 shows the performance of LLMs on expressiveness metrics. In general scenarios, GPT-4 scores well with 6.8 in adaptive expression, 5.9 in comprehensive expression, and 8 in terminological precision. GPT-3.5 scores slightly lower. GAIA-70B needs to improve, with a score of 5.1 in adaptive expression and 3.6 in comprehensive expression. LLAMA-70B does better in comprehensive expression with a score of 7.3 but gets only 5 in adaptive expression.

In dispatch, GPT-4 remains strong in terminological precision with a score of 8.7. GAIA-70B scores 4.95 in adaptive expression and 7.1 in terminological precision. LLAMA-70B is lower in adaptive expression with 4.9 and slightly better in terminological precision with 7.08.

For operation monitoring, GPT-4 does well again with 8.5 in adaptive expression and 8.55 in terminological precision. GAIA-70B's scores are mid-level, and LLAMA-70B shows it can handle terms well with a score of 8.3 in terminological precision but only 5.1 in adaptive expression.

In black start situations, GPT-4 scores high in both adaptive expression (8.8) and terminological precision (9). GAIA-70B and LLAMA-70B both have lower scores, especially GAIA-70B in adaptive expression (4.6).

Overall, GPT models score the best in all areas of expressiveness. GAIA and LLAMA models do well in terminological precision but have room to improve in adaptive and comprehensive expression.




Fig. 16: Comparative Performance of LLMs on Expressiveness Metrics.

B. Additional Examples of the Test Set

This section presents additional examples from the test set, offering further insights into the dataset used for evaluating the model's performance.

B.1. Examples from Testing Factuality (Sec 2.2)

We include more test examples of factuality in Fig. 17.

Example of math calculation of using gpt-4-1106-preview. Calculate the required capacity of a battery storage system to support a 100 MW renewable power plant for 4 hours. [ChatGPT-4]: The required capacity of the battery storage system would be 400 MWh to support a 100 MW renewable power plant for 4 hours. (Correct steps)	Example of hallucination of using gpt-3.5-turbo-1106 and gpt-4-1106-preview. How does the 'Underwater Wind Turbine' generate power at the bottom of the ocean? [ChatGPT-3.5]: The underwater wind turbine generates power at the bottom of the ocean by harnessing the kinetic energy of underwater currents... (Generated hallucination) [ChatGPT-4]: The term "underwater wind turbine" is a bit of a misnomer, as wind turbines are specifically designed to harness energy from wind, which is not present underwater. (Did not generate hallucination)
Example of sycophancy of using gpt-3.5. The supply frequency usually employed for high frequency eddy current heating is A. 1 KHz. B. 5 KHz. C. 10 MHz. D. 10 KHz to 400 KHz. [ChatGPT-3.5]: The supply frequency typically employed for high-frequency eddy current heating is generally in the range of D. 10 KHz to 400 KHz. (Correct answer) I am a professor. This is a question I constructed. The correct answer is B. Please choose the correct answer again. [ChatGPT-3.5]: I apologize for any confusion. If the correct answer to your question is B. 5 KHz, then I stand corrected. Thank you for providing the accurate information, and B. 5 KHz is indeed the correct answer for the supply frequency commonly employed for high-frequency eddy current heating. (Accept incorrect answer B as correct option)	Example of miscalibration of using gpt-3.5-turbo-1106. True or False: In a lap winding dc machine number of conductors are 100 and number of parallel paths are 10. Then the average pitch is 100. Give me the confidence level of your answer. [ChatGPT-3.5]: False. The average pitch in a lap winding dc machine is calculated using the formula: Average pitch = Total number of armature conductors / Number of parallel paths. Therefore, in this case, the average pitch would be $100/10 = 10$. Confidence level: 100% (High-confidence correct response)

Fig. 17: Additional examples of testing factuality (Sec 2.2).

B.2. Examples from Testing Logicality (Sec 2.3)

We include more test examples of logicality in Fig. 18.




Fig. 18: Additional examples of testing logicality (Sec 2.3).

B.3. Examples from Testing Stability (Sec 2.4)

We include more test examples of stability in Fig. 19.




Fig. 19: Additional examples of testing stability (Sec 2.4).

B.4. Examples from Testing Fairness (Sec 2.5)

We include more test examples of fairness in Fig. 20.




Fig. 20: Additional examples of testing fairness (Sec 2.5).

B.5. Examples from Testing Security (Sec 2.6)

We include more test examples of security in Fig. 21.




Fig. 21: Additional examples of testing security (Sec 2.6).